

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
import statsmodels.api as sm
from scipy.stats import f_oneway, ttest_ind, chi2_contingency
from mpl_toolkits.mplot3d import Axes3D

cars_df = pd.read_csv('/content/drive/MyDrive/5502/car_data_set.csv')
```

```
cars_df.shape

(11914, 16)
```

```
cars_df.head(5)
```

	Make	Model	Year	Engine Fuel Type	Engine HP	Engine Cylinders	Transmission Type	Driven_Wheels	Number of Doors	Market	Category	Vehicle Size	Vehicle Style	highway MPG	c
0	BMW	Series 1 M	2011	premium unleaded (required)	335.0	6.0	MANUAL	rear wheel drive	2.0	Tuner,Factory,Luxury,High-Performance		Compact	Coupe	26	
1	BMW	Series 1	2011	premium unleaded (required)	300.0	6.0	MANUAL	rear wheel drive	2.0	Luxury,Performance		Compact	Convertible	28	
2	BMW	Series 1	2011	premium unleaded (required)	300.0	6.0	MANUAL	rear wheel drive	2.0	Luxury,High-Performance		Compact	Coupe	28	
3	BMW	Series 1	2011	premium unleaded (required)	230.0	6.0	MANUAL	rear wheel drive	2.0	Luxury,Performance		Compact	Coupe	28	
4	BMW	Series 1	2011	premium unleaded (required)	230.0	6.0	MANUAL	rear wheel drive	2.0		Luxury	Compact	Convertible	28	

```
print(cars_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11914 entries, 0 to 11913
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Make                   11914 non-null object
1   Model                  11914 non-null object
2   Year                   11914 non-null int64
3   Engine Fuel Type      11911 non-null object
4   Engine HP              11845 non-null float64
5   Engine Cylinders      11884 non-null float64
6   Transmission Type     11914 non-null object
7   Driven_Wheels         11914 non-null object
8   Number of Doors       11908 non-null float64
9   Market Category       8172 non-null  object
10  Vehicle Size           11914 non-null object
11  Vehicle Style          11914 non-null object
12  highway MPG            11914 non-null int64
13  city mpg               11914 non-null int64
14  Popularity             11914 non-null int64
15  MSRP                   11914 non-null int64
dtypes: float64(3), int64(5), object(8)
memory usage: 1.5+ MB
None
```

```
print(cars_df.dtypes)
```

```
Make           object
Model          object
Year           int64
Engine Fuel Type  object
Engine HP      float64
Engine Cylinders float64
```

```
Transmission Type    object
Driven_Wheels        object
Number of Doors      float64
Market Category      object
Vehicle Size         object
Vehicle Style        object
highway MPG          int64
city mpg             int64
Popularity           int64
MSRP                 int64
dtype: object
```

```
cars_df.describe()
```

	Year	Engine HP	Engine Cylinders	Number of Doors	highway MPG	city mpg	Popularity	MSRP
count	11914.000000	11845.00000	11884.000000	11908.000000	11914.000000	11914.000000	11914.000000	1.191400e+04
mean	2010.384338	249.38607	5.628829	3.436093	26.637485	19.733255	1554.911197	4.059474e+04
std	7.579740	109.19187	1.780559	0.881315	8.863001	8.987798	1441.855347	6.010910e+04
min	1990.000000	55.00000	0.000000	2.000000	12.000000	7.000000	2.000000	2.000000e+03
25%	2007.000000	170.00000	4.000000	2.000000	22.000000	16.000000	549.000000	2.100000e+04
50%	2015.000000	227.00000	6.000000	4.000000	26.000000	18.000000	1385.000000	2.999500e+04
75%	2016.000000	300.00000	6.000000	4.000000	30.000000	22.000000	2009.000000	4.223125e+04
max	2017.000000	1001.00000	16.000000	4.000000	354.000000	137.000000	5657.000000	2.065902e+06

```
def check_missing_values(df):
    null_values = df.isnull().sum()
    print("Null Values:")
    print(null_values)

    nan_values = df.isna().sum()
    print("\nNaN Values:")
    print(nan_values)

    empty_values = (df == '').sum()
    print("\nEmpty Values:")
    print(empty_values)

def preprocess_data(df, numeric_columns, categorical_columns):
    for col in numeric_columns:
        col_mean = df[col].mean()
        df[col].fillna(col_mean, inplace=True)

    for col in categorical_columns:
        df[col].fillna(df[col].mode()[0], inplace=True)

    return df

check_missing_values(cars_df)

numeric_cols = ['Engine HP', 'Engine Cylinders', 'Number of Doors', 'highway MPG', 'city mpg', 'Popularity', 'MSRP']
categorical_cols = ['Engine Fuel Type', 'Market Category']

cars_df = preprocess_data(cars_df, numeric_cols, categorical_cols)

check_missing_values(cars_df)
```

```
vehicle Size      0
Vehicle Style     0
highway MPG       0
city mpg          0
Popularity        0
MSRP              0
dtype: int64
```

NaN Values:

```
Make      0
Model     0
Year      0
Engine Fuel Type  0
Engine HP  0
Engine Cylinders  0
Transmission Type  0
Driven_Wheels  0
Number of Doors  0
Market Category  0
Vehicle Size  0
Vehicle Style  0
highway MPG   0
city mpg      0
Popularity    0
MSRP          0
dtype: int64
```


Empty Values:

```
Make      0
Model     0
Year      0
Engine Fuel Type  0
Engine HP  0
Engine Cylinders  0
Transmission Type  0
Driven_Wheels  0
Number of Doors  0
Market Category  0
Vehicle Size  0
Vehicle Style  0
highway MPG   0
city mpg      0
Popularity    0
MSRP          0
dtype: int64
```

CORRELATION HEATMAP


```
plt.figure(figsize=(10, 8))
sns.heatmap(cars_df.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
```

```
<ipython-input-8-17d755946109>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version,
sns.heatmap(cars_df.corr(), annot=True, cmap='coolwarm', fmt='.2f')
```



Numerical vs Numerical

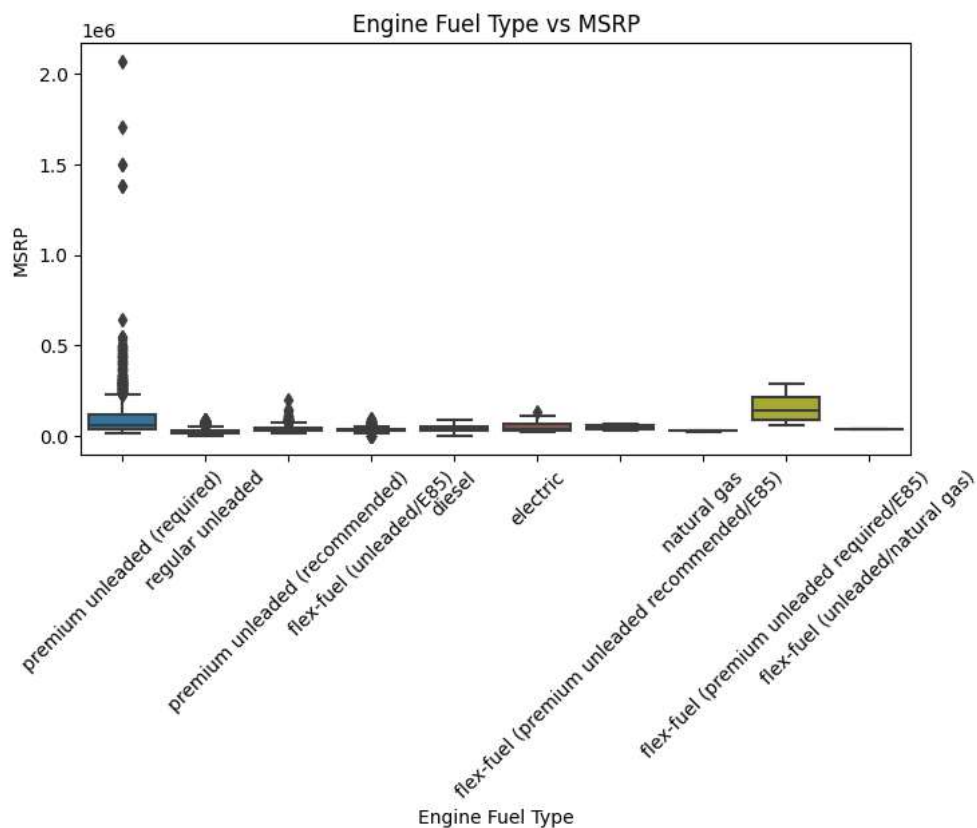




Numerical vs Categorical:

Boxplot: Engine Fuel Type vs MSRP


```
plt.figure(figsize=(8, 4))
sns.boxplot(x='Engine Fuel Type', y='MSRP', data=cars_df)
plt.title('Engine Fuel Type vs MSRP')
plt.xticks(rotation=45)
plt.show()
```



Numerical vs Categorical:

Barplot: Transmission Type vs MSRP


```
plt.figure(figsize=(6, 6))
sns.barplot(x='Transmission Type', y='MSRP', data=cars_df)
plt.title('Transmission Type vs MSRP')
plt.xticks(rotation=45)
plt.show()
```



Categorical vs Categorical:

Countplot: Vehicle Style by Drivetrain


```
plt.figure(figsize=(8, 6))
sns.countplot(x='Vehicle Style', hue='Driven_Wheels', data=cars_df)
plt.title('Vehicle Style by Drivetrain')
plt.legend(loc='upper right')
plt.xticks(rotation=45)
plt.show()
```



Numerical and Categorical vs Categorical:

Barplot: Engine Cylinders & Vehicle Size by Market Category


```
plt.figure(figsize=(12, 16))
sns.barplot(x='Engine Cylinders', y='Market Category', hue='Vehicle Size', data=cars_df)
plt.title('Engine Cylinders & Vehicle Size by Market Category')
plt.legend(loc='upper right')
plt.show()
```



Two Numerical vs One Categorical:

Scatterplot: Engine HP, Highway MPG by Vehicle Style

```
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Engine HP', y='highway MPG', hue='Transmission Type', data=cars_df)
plt.title('Engine HP, Highway MPG by Vehicle Style')
plt.legend(loc='upper right')
plt.show()
```



Three Numerical Variables:

3D Scatterplot: Engine HP, Highway MPG, City MPG

```
Q1 = cars_df.quantile(0.25)
Q3 = cars_df.quantile(0.75)
IQR = Q3 - Q1

outliers = ((cars_df < (Q1 - 1.5 * IQR)) | (cars_df > (Q3 + 1.5 * IQR))).sum()
print("Potential Outliers (using IQR method):")
print(outliers)
```

Potential Outliers (using IQR method):

Driven_Wheels	0
Engine_Cylinders	357
Engine_Fuel_Type	0
Engine_HP	509
MSRP	996
Make	0
Market_Category	0
Model	0
Number_of_Doors	0
Popularity	881
Transmission_Type	0
Vehicle_Size	0
Vehicle_Style	0
Year	661
city_mpg	316
highway_MPG	192

dtype: int64

<ipython-input-15-73b63a5a9663>:1: FutureWarning: The default value of numeric_only in DataFrame.quantile is deprecated. In a future ver


```

Q1 = cars_df.quantile(0.25)
<ipython-input-15-73b63a5a9663>:2: FutureWarning: The default value of numeric_only in DataFrame.quantile is deprecated. In a future ver
Q3 = cars_df.quantile(0.75)
<ipython-input-15-73b63a5a9663>:5: FutureWarning: Automatic reindexing on DataFrame vs Series comparisons is deprecated and will raise v
outliers = ((cars_df < (Q1 - 1.5 * IQR)) | (cars_df > (Q3 + 1.5 * IQR))).sum()

```

T-TEST

```

premium_unleaded = cars_df[cars_df['Engine Fuel Type'] == 'premium unleaded (required)']
regular_unleaded = cars_df[cars_df['Engine Fuel Type'] == 'regular unleaded']

t_stat, p_value = stats.ttest_ind(premium_unleaded['MSRP'], regular_unleaded['MSRP'])
print(f"T-test results - t-statistic: {t_stat}, p-value: {p_value}")

alpha = 0.05 #significance level

if p_value < alpha:
    print("Reject null hypothesis: There is a significant difference in MSRP based on engine type.")
else:
    print("Fail to reject null hypothesis: No significant difference in MSRP based on engine type.")

T-test results - t-statistic: 53.05815695600708, p-value: 0.0
Reject null hypothesis: There is a significant difference in MSRP based on engine type.

```

Correlation Analysis

```

correlation_matrix = cars_df[['Engine HP', 'MSRP']].corr(method='pearson')
print("Correlation between Engine HP and MSRP:")
print(correlation_matrix)

```

```

Correlation between Engine HP and MSRP:
Engine HP      MSRP
Engine HP    1.000000  0.661644
MSRP         0.661644  1.000000

```

ANOVA TEST

```

grouped_data = cars_df.groupby('Vehicle Style')['MSRP'].apply(list)
f_statistic, p_value = f_oneway(*grouped_data)
print("\nANOVA Test - MSRP variation among Vehicle Styles:")
print(f"F-Statistic: {f_statistic}, p-value: {p_value}")

```

```

ANOVA Test - MSRP variation among Vehicle Styles:
F-Statistic: 92.55549649004786, p-value: 5.613983147064976e-271

```

T - TEST

```

engine_fuel_types = cars_df['Engine Fuel Type'].unique()

for fuel_type in engine_fuel_types:
    subset = cars_df[cars_df['Engine Fuel Type'] == fuel_type]['MSRP']
    t_stat, p_value = ttest_ind(subset, cars_df['MSRP'])
    print(f"\nT-test for {fuel_type} Engine Fuel Type and MSRP:")
    print(f"T-statistic: {t_stat}, p-value: {p_value}")

```

```

T-test for premium unleaded (required) Engine Fuel Type and MSRP:
T-statistic: 34.910868116537976, p-value: 4.7306254324863656e-256

```

T-test for regular unleaded Engine Fuel Type and MSRP:
T-statistic: -24.345905204481735, p-value: 5.905242990090284e-129

T-test for premium unleaded (recommended) Engine Fuel Type and MSRP:
T-statistic: 0.14094797183948202, p-value: 0.8879131574803788

T-test for flex-fuel (unleaded/E85) Engine Fuel Type and MSRP:
T-statistic: -2.158867095835616, p-value: 0.030878987504420384

T-test for diesel Engine Fuel Type and MSRP:
T-statistic: 0.039873987036859695, p-value: 0.9681942496423356

T-test for electric Engine Fuel Type and MSRP:
T-statistic: 0.9926277529140373, p-value: 0.3209114461704788

T-test for flex-fuel (premium unleaded recommended/E85) Engine Fuel Type and MSRP:
T-statistic: 0.6825686452754542, p-value: 0.4948926966706487

T-test for natural gas Engine Fuel Type and MSRP:
T-statistic: -0.2947803041303683, p-value: 0.7681668493734748

T-test for flex-fuel (premium unleaded required/E85) Engine Fuel Type and MSRP:
T-statistic: 14.483807998706142, p-value: 3.838640790513522e-47

T-test for flex-fuel (unleaded/natural gas) Engine Fuel Type and MSRP:
T-statistic: -0.05707185895069415, p-value: 0.9544889099295308

Chi-square Test:

```
cross_tab = pd.crosstab(cars_df['Vehicle Style'], cars_df['Market Category'])
chi2, p, dof, expected = chi2_contingency(cross_tab)
print("\nChi-square Test - Association between Vehicle Style and Market Category:")
print(f"Chi-square value: {chi2}, p-value: {p}")
```

Chi-square Test - Association between Vehicle Style and Market Category:
Chi-square value: 27100.05046976974, p-value: 0.0

```
X = cars_df['Engine HP'] # Independent variable
y = cars_df['MSRP'] # Dependent variable

X = sm.add_constant(X) # Add a constant term for intercept
model = sm.OLS(y, X).fit() # Fit the model
print(model.summary())
```

OLS Regression Results						
Dep. Variable:	MSRP	R-squared:	0.438			
Model:	OLS	Adj. R-squared:	0.438			
Method:	Least Squares	F-statistic:	9275.			
Date:	Thu, 16 Nov 2023	Prob (F-statistic):	0.00			
Time:	23:54:11	Log-Likelihood:	-1.4458e+05			
No. Observations:	11914	AIC:	2.892e+05			
Df Residuals:	11912	BIC:	2.892e+05			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-5.05e+04	1032.111	-48.932	0.000	-5.25e+04	-4.85e+04
Engine HP	365.2884	3.793	96.308	0.000	357.854	372.723
Omnibus:	23521.060	Durbin-Watson:	0.713			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	105879175.386			
Skew:	15.628	Prob(JB):	0.00			
Kurtosis:	463.771	Cond. No.	680.			

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

